



# Oracle-guided Dynamic User Preference Modeling for Sequential Recommendation

Jiafeng Xia\*  
Fudan University  
Shanghai, China  
jfxia19@fudan.edu.cn

Dongsheng Li  
Microsoft Research Asia  
Shanghai, China  
dongсли@microsoft.com

Hansu Gu  
Amazon  
Seattle, United States  
hansug@acm.org

Tun Lu\*†‡  
Fudan University  
Shanghai, China  
lutun@fudan.edu.cn

Peng Zhang\*‡  
Fudan University  
Shanghai, China  
zhangpeng\_@fudan.edu.cn

Li Shang\*  
Fudan University  
Shanghai, China  
lishang@fudan.edu.cn

Ning Gu\*  
Fudan University  
Shanghai, China  
ninggu@fudan.edu.cn

## ABSTRACT

Sequential recommendation methods can capture dynamic user preferences from user historical interactions to achieve better performance. However, most existing methods only use past information extracted from user historical interactions to train the models, leading to the deviations of user preference modeling. Besides past information, future information is also available during training, which contains the “oracle” user preferences in the future and will be beneficial to model dynamic user preferences. Therefore, we propose an oracle-guided dynamic user preference modeling method for sequential recommendation (Oracle4Rec), which leverages future information to guide model training on past information, aiming to learn “forward-looking” models. Specifically, Oracle4Rec first extracts past and future information through two separate encoders, then learns a forward-looking model through an oracle-guiding module which minimizes the discrepancy between past and future information. We also tailor a two-phase model training strategy to make the guiding more effective. Extensive experiments demonstrate that Oracle4Rec is superior to state-of-the-art sequential methods. Further experiments show that Oracle4Rec can be leveraged as a generic module in other sequential recommendation methods to improve their performance with a considerable margin.

\*Also with School of Computer Science and Shanghai Key Laboratory of Data Science, Fudan University.

†Also with Fudan Institute on Aging, MOE Laboratory for National Development and Intelligent Governance, and Shanghai Institute of Intelligent Electronics & Systems, Fudan University.

‡Corresponding author.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](https://permissions.acm.org).

WSDM '25, March 10–14, 2025, Hannover, Germany

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 979-8-4007-1329-3/25/03

<https://doi.org/10.1145/3701551.3703542>

## CCS CONCEPTS

• Information systems → Recommender systems.

## KEYWORDS

Dynamic user preference modeling, sequential recommendation

### ACM Reference Format:

Jiafeng Xia, Dongsheng Li, Hansu Gu, Tun Lu, Peng Zhang, Li Shang, and Ning Gu. 2025. Oracle-guided Dynamic User Preference Modeling for Sequential Recommendation. In *Proceedings of the Eighteenth ACM International Conference on Web Search and Data Mining (WSDM '25)*, March 10–14, 2025, Hannover, Germany. ACM, New York, NY, USA, 10 pages. <https://doi.org/10.1145/3701551.3703542>

## 1 INTRODUCTION

Recommender systems, which can recommend potentially interested items to the users according to their historical interactions, have been widely used in various fields, e.g., advertising [24, 28], movie recommendation [12, 40] and E-commerce [3, 30, 35]. Generally, user interactions are continuous and can be seen as a sequence of user’s interacted items sorted in chronological order, which makes user preferences constantly change over time in nature. Sequential recommendation methods [21, 37, 46, 47, 49, 55] can model dynamic user preferences based on his/her historical interaction sequence and achieve better performance compared with static recommendation algorithms [4, 5, 27, 34, 48].

However, existing sequential recommendation methods are faced with one key challenge. When predicting the next interaction, they only use past information extracted from user interactions that come before the current interaction, which could cause the deviations of dynamic user preference modeling since only using past information is insufficient to capture the dynamics of user preferences [53], and thus hurt the performance of recommendation models. Besides past information, future information, which can be extracted from user interactions that come after the current interaction, is also available in the training. It can be viewed as a type of posterior information containing the “oracle” user preferences in the future and is also beneficial to model dynamic user preferences.

Using future information to guide model training on past information can learn forward-looking models, which can better model dynamic user preference and thus improve model performance.

To this end, we propose an oracle-guided dynamic user preference modeling method for sequential recommendation (Oracle4Rec) that leverages future information to guide model training on past information, aiming to learn forward-looking models. Specifically, Oracle4Rec first extracts past information and future information through two separate information encoders respectively. The two encoders have the same architecture, which both have a noise filtering module to filter the noise in the user interaction sequence, followed by a causal self-attention module to capture the evolution of user preferences, making the extracted information more accurate. Then Oracle4Rec learns the forward-looking model through a carefully designed oracle-guiding module in a manner of minimizing the discrepancy between past information and future information. To make the guiding from future information to past information more effective, we also tailor a two-phase model training strategy, named 2PTraining, for Oracle4Rec. Extensive experiments on six real-world datasets demonstrate that Oracle4Rec consistently outperforms state-of-the-art sequential recommendation algorithms. We further implement Oracle4Rec as a generic module due to it is orthogonal to existing sequential methods, and we conduct a generality experiment to show that Oracle4Rec has high generality and it can be flexibly applied to other sequential methods and improve their performance with a considerable margin.

The contributions of this work are summarized as follows:

- We propose an oracle-guided dynamic user preference modeling method for sequential recommendation that can learn forward-looking models through a carefully designed oracle-guiding module and a tailored training strategy, so as to better model dynamic user preferences.
- We conduct extensive experiments on six public datasets, and the results show that Oracle4Rec achieves better performance compared with other state-of-the-art sequential methods.
- We implement Oracle4Rec<sup>1</sup> as a generic module due to it is orthogonal to existing sequential methods, and experiments show that it can improve the performance of those methods with a considerable margin.

## 2 THE PROPOSED METHOD

### 2.1 Notations

Let user set and item set be  $\mathcal{U}$  and  $\mathcal{V}$  respectively,  $|\mathcal{U}| = m$  and  $|\mathcal{V}| = n$ . The user  $u$ 's interaction sequence in chronological order can be denoted as  $\mathcal{S}^u = \{v_i | v_i \in \mathcal{V}, i = 1, 2, \dots, |\mathcal{S}^u|\}$ . Our goal is to predict next item  $v_{|\mathcal{S}^u|+1}$  that user  $u$  will interact with based on his/her interaction sequence  $\mathcal{S}^u$ . For the convenience of the following description, we give the definitions of two kinds of user interaction sequences. For a target interaction  $v_t \in \mathcal{V}$  that needs to be predicted for user  $u$ , his/her historical interaction sequence can be defined as  $\mathcal{H}_t^u = \{v_1, v_2, \dots, v_{t-1}\}$ , and his/her global interaction sequence is defined as  $\mathcal{G}_t^u = \{v_1, v_2, \dots, v_{t-1}, v_t, v_{t+1}, \dots, v_{t+P}\}$ , which additionally contains  $P + 1$  items compared with user historical interaction sequence. We usually adopt fixed-length historical

interaction sequence and global interaction sequence for high efficiency during training and inference, thus we need to truncate or pad the sequences  $\mathcal{H}_t^u$  and  $\mathcal{G}_t^u$  to ensure their lengths are equal to a positive integer  $L (L > P)$ . If the length of sequence is longer than  $L$ , we truncate its length to  $L$  by selecting the recent  $L$  interactions, while if the length of sequence is shorter than  $L$ , we repeatedly add a padding item to the left side of sequence until its length is  $L$ . Therefore, we redefine the historical interaction sequence as  $\mathcal{H}_t^u = \{v_{t-L}, v_{t-L+1}, \dots, v_{t-1}\}$  and the global interaction sequence as  $\mathcal{G}_t^u = \{v_{t+P-L+1}, v_{t+P-L+2}, \dots, v_{t+P}\}$ .

### 2.2 Overview of Oracle4Rec

Oracle4Rec is a sequential recommendation method that leverages past information and future information to jointly model dynamic user preferences and make recommendation. Figure 1 shows the architecture of the proposed Oracle4Rec, which consists of three key components: a **Past Information Encoder** used to extract past information from user historical interaction sequences, a **Future Information Encoder** used to extract future information from user global interaction sequences, and an **Oracle-Guiding Module** used to guide model training on past information with future information. We also tailor an effective model training strategy 2PTraining to facilitate the guiding from future information to past information, thereby making past information capture the dynamics of user preferences more accurately and sufficiently. Next, we introduce these modules and the tailored training strategy in details.

### 2.3 Past Information Encoder

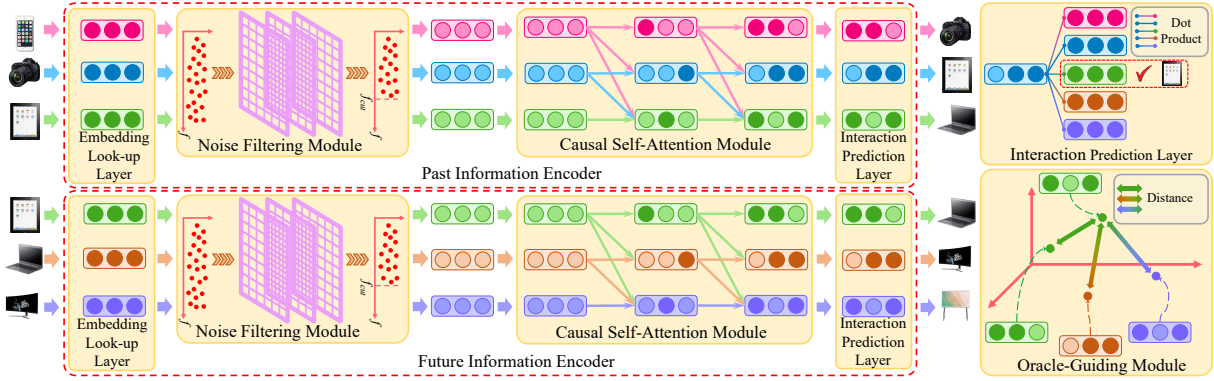
**2.3.1 Embedding Look-up Layer.** Embedding Look-up Layer is used to generate item embedding matrix  $\mathbf{E}^u \in \mathbb{R}^{L \times d}$  according to the user historical interaction sequence  $\mathcal{H}_t^u$  by looking up the item embedding table  $\mathbf{T} \in \mathbb{R}^{|\mathcal{V}| \times d}$ . However, unlike GRU [6] and LSTM [18], which can capture the order relation between items in the sequence, the attention mechanism in the following Causal Self-Attention Module will ignore it, affecting the accuracy of the dynamic user preference modeling. Therefore, we introduce the positional embedding  $\mathbf{E}^P \in \mathbb{R}^{L \times d}$  to help the model to identify the order relation between different items. There are many choices of positional embedding, such as sinusoidal positional embedding [42] and learnable positional embedding [7], we follow the suggestion in [21] to use learnable positional embedding for better performance. Thus, the final item embedding is:

$$\hat{\mathbf{E}}^u = \mathbf{E}^u + \mathbf{E}^P. \quad (1)$$

Note that we will take user  $u$ 's historical sequence  $\mathcal{H}_t^u$  as an example to illustrate how past information encoder works. Thus, we will omit the superscript of all embeddings that are associated with user  $u$  for simplicity hereinafter.

**2.3.2 Noise Filtering Module.** User interaction sequence will inevitably introduce some noisy interactions due to the first impression [43], such as caption bias [19] and position bias [20], which will cause deviations in modeling user preferences. Specifically, the self-attention mechanism employed in the Causal Self-Attention Module (Section 2.3.3) is highly sensitive to noisy interactions, and the distribution of attention within the model is easily influenced

<sup>1</sup>We release the code for further research: <https://github.com/Yaveng/Oracle4Rec>.



**Figure 1: The architecture of Oracle4Rec, which consists of three parts: Past Information Encoder, Future Information Encoder and Oracle-Guiding Module. The first two encoders are both composed of Embedding Look-up Layer, Noise Filtering Module, Causal Self-Attention Module and Interaction Prediction Layer. For the ease of presentation, the Oracle-Guiding Module is set to minimize the distance between past information and future information in a 3D coordinate system.**

by these noisy interactions, resulting in certain biases in the extracted information. To mitigate the impact of noisy interactions on modeling user dynamic preference, we introduce a Noise Filtering Module before the Causal Self-Attention Module. This adjustment reduces the influence of noisy interactions on modeling user dynamic preference, thereby improving the accuracy of user future interaction predictions.

In the signal processing field, researchers usually analyze and reduce the noise of time series signal in the frequency domain [8]. Specifically, the time series signal is first transformed from time domain to frequency domain through Fourier transform, so as to obtain a series of frequency components. Generally, noise corresponds to high-frequency components [52]. By eliminating a certain proportion of high-frequency components, we can achieve noise reduction of the signal. Then we transform the signal to time domain through inverse Fourier transform, and we obtain a noise-free time series signal. FMLP-Rec [55] designs a learnable filter to attenuate the noise in the item embedding by imitating the above process. However, the learnable filter may introduce additional model parameters and increase the difficulty of model training.

Since noise corresponds to the high-frequency components and our goal is to remove the noise in the item embedding. Thus, we design the Noise Filtering Module to remove high-frequency components in the item embedding to achieve noise reduction. The noise filtering module has multiple layers, each layer contains a parameter-free low-pass filter to specify how many low-frequency components are preserved. Concretely, in the  $g$ -th layer, we first transform the item embedding  $\mathbf{F}^{g-1}$  from time domain to the frequency domain through the Fast Fourier Transform (FFT):

$$\mathbf{X}^g, \mathbf{f}^g \leftarrow \text{FFT}(\mathbf{F}^{g-1}), \quad (2)$$

where  $\mathbf{F}^{g-1}$  is the item embedding from the  $(g-1)$ -th layer and  $\mathbf{F}^0 = \hat{\mathbf{E}}$ ,  $\mathbf{X}^g \in \mathbb{C}^{c \times d}$  and  $\mathbf{f}^g \in \mathbb{R}^c$  are the spectrum matrix and frequency vector respectively,  $c$  is the number of frequency components. The cutoff frequency  $f_{cut}^g$  can be obtained by calculating the  $q$  quantile of  $\mathbf{f}^g$ . Then the indicator matrix  $\mathbf{M}^g \in \mathbb{R}^{c \times d}$ , which decides whether the frequency components are preserved or removed,

can be calculated as:

$$\mathbf{M}^g = \underbrace{\left[ \mathbb{1}_{f_1^g < f_{cut}^g}, \mathbb{1}_{f_2^g < f_{cut}^g}, \dots, \mathbb{1}_{f_c^g < f_{cut}^g} \right]^T}_{\text{contains } c \text{ elements}} \otimes \mathbf{1}^T, \quad (3)$$

where  $\mathbb{1}$  is an indicator function,  $\mathbf{1}$  is a  $d$ -dimensional vector whose elements are all 1 and  $\otimes$  is the outer product. The filtered spectrum  $\tilde{\mathbf{X}}^g$  can be obtained by an element-wise product of  $\mathbf{X}^g$  and  $\mathbf{M}^g$ :

$$\tilde{\mathbf{X}}^g = \mathbf{X}^g \odot \mathbf{M}^g. \quad (4)$$

Then we can get the item embedding after noise reduction through the Inverse Fast Fourier Transform (IFFT). Meanwhile, we incorporate the skip connection [14], layer normalization [2] and dropout [36] to stabilize training. Thus, the final item embedding  $\mathbf{F}^g$  is

$$\mathbf{F}^g = \text{LayerNorm}(\text{Dropout}(\text{IFFT}(\tilde{\mathbf{X}}^g)) + \mathbf{F}^{g-1}). \quad (5)$$

After  $G$  layers, the output of noise filtering module is  $\mathbf{F}$ .

**2.3.3 Causal Self-Attention Module.** Compared with RNNs, attention mechanism has surprising long-range dependency modeling ability [11, 38]. Therefore, we use the self-attention mechanism, which borrows from [21], to construct the Causal Self-Attention Module, so as to capture the dynamic user preferences after reducing the noise in the item embedding. The causal self-attention module is also composed of multiple layers. In the  $k$ -th layer, the item embedding  $\mathbf{S}^{k-1}$  is first fed into a self-attention layer to capture the global information for all items in the sequence:

$$\begin{aligned} \mathbf{Z}^k &= \text{Softmax}((\mathbf{S}^{k-1} \mathbf{W}_Q^k)(\mathbf{S}^{k-1} \mathbf{W}_K^k)^T / \sqrt{d}), \\ \tilde{\mathbf{O}}^k &= \mathbf{Z}^k (\mathbf{S}^{k-1} \mathbf{W}_V^k), \\ \mathbf{O}^k &= \text{LayerNorm}(\text{Dropout}(\tilde{\mathbf{O}}^k \mathbf{W}_1^k + \mathbf{b}_1^k) + \mathbf{S}^{k-1}), \end{aligned} \quad (6)$$

where  $\mathbf{W}_Q^k$ ,  $\mathbf{W}_K^k$  and  $\mathbf{W}_V^k$  are queries, keys and values projection matrices, and  $\mathbf{W}_1^k$  and  $\mathbf{b}_1^k$  are weight and bias of linear transformation at  $k$ -th layer. Skip connection, layer normalization and dropout are used to stabilize training.  $\mathbf{S}^{k-1}$  is the item embedding in the  $(k-1)$ -th layer and  $\mathbf{S}^0 = \mathbf{F}$ . Note that self-attention mechanism does not guarantee the causality that requires the former items cannot

receive information from the latter items. Violating the causality will prevent model from learning useful feature and have bad performance in inference. Therefore, when calculating  $\mathbf{Z}^k$  in Eq. (6), we should set  $\mathbf{Z}_{ij}^k$  to 0 if  $i < j$  ( $j = 1, 2, \dots, n$ ).

Then we apply a feed-forward network to the item embedding  $\mathbf{O}^k$  to capture the non-linearity in the embedding:

$$\mathbf{H}^k = \sigma(\mathbf{O}^k \mathbf{W}_2^k + \mathbf{b}_2^k) \mathbf{W}_3^k + \mathbf{b}_3^k, \quad (7)$$

where  $\mathbf{W}_2^k$ ,  $\mathbf{W}_3^k$ ,  $\mathbf{b}_2^k$  and  $\mathbf{b}_3^k$  are weights and biases,  $\sigma(\cdot)$  is the activation function. We also incorporate the skip connection, layer normalization and dropout to stabilize training:

$$\mathbf{S}^k = \text{LayerNorm}(\text{Dropout}(\mathbf{H}^k) + \mathbf{O}^k). \quad (8)$$

After  $K$  layers, we denote the output of this module as  $\mathbf{S}$ .

**2.3.4 Interaction Prediction Layer.** After obtaining the predicted item embedding  $\mathbf{S}$ , we can calculate the probability  $P(v_t | \mathcal{H}_t^u)$  of the target interaction  $v_t$  based on  $\mathcal{H}_t^u$  as:

$$P(v_t | \mathcal{H}_t^u) = \text{Sigmoid}(\mathbf{Q}_L^T \mathbf{T}_{v_t}), \quad (9)$$

where  $\mathbf{Q}_L$  is the predicted embedding of the target interaction, which corresponds to the last row of  $\mathbf{Q}$ ,  $\mathbf{T}_{v_t}$  is the real embedding of the target interaction from embedding table  $\mathbf{T}$ . Therefore, the loss function of past information encoder is:

$$\mathcal{L}_p = - \sum_{u=1}^{|U|} \sum_{t=1}^{|S^u|} [\log(P(v_t | \mathcal{H}_t^u)) + \log(1 - P(j | \mathcal{H}_t^u))], \quad (10)$$

where  $S^u$  is user  $u$ 's interaction sequence, and  $j \notin S^u$  is a negative item that user  $u$  has never interacted with.

## 2.4 Future Information Encoder

Future information encoder is used to encode the user future preferences from user global interaction sequence. It has the same structure as the past information encoder, which consists of embedding look-up layer, noise filtering module, causal self-attention module and interaction prediction layer. Different from BERT4Rec [37], GRec [53] and DualRec [54] that adopt the reversed user interaction sequence to extract future information, i.e., in a Right-to-Left-style, our future information encoder extracts future information in a more natural way, i.e., in a Left-to-Right-style. Given user  $u$ 's global interaction sequence  $\mathcal{G}_t^u$ , the future information encoder outputs the item embedding  $\mathbf{R}$ . Then we can calculate the probability  $P(v_{t+P+1} | \mathcal{G}_t^u)$  of the interaction  $v_{t+P+1}$ :

$$P(v_{t+P+1} | \mathcal{G}_t^u) = \text{Sigmoid}(\mathbf{R}_L^T \mathbf{T}_{v_{t+P+1}}), \quad (11)$$

where  $\mathbf{R}_L$  is the predicted embedding of item  $v_{t+P+1}$  and it is the last row of  $\mathbf{R}$ ,  $\mathbf{T}_{v_{t+P+1}}$  is the real embedding of item  $v_{t+P+1}$  from embedding table  $\mathbf{T}$ , and  $P$  is the number of future interactions after the target interaction  $v_t$  in  $\mathcal{G}_t^u$ . However, the loss function of future information encoder is slightly different from that of past information encoder. Since we want to incorporate future information to make user preference modeling more accurate, we should ensure that future information is sufficiently extracted from user global interaction sequence. Therefore, we use the following

loss function to optimize the future information encoder:

$$\begin{aligned} \mathcal{L}_f &= - \sum_{u=1}^{|U|} \sum_{t=1}^{|S^u|} [\log(P(\mathcal{G}_{t+1}^u | \mathcal{G}_t^u)) + \log(1 - P(\mathcal{J} | \mathcal{G}_t^u))] \\ &= - \sum_{u=1}^{|U|} \sum_{t=1}^{|S^u|} \sum_{l=1}^L [\log(P(v_{m+l} | \mathcal{G}_{t,l}^u)) + \log(1 - P(j_l | \mathcal{G}_{t,l}^u))], \end{aligned} \quad (12)$$

where  $\mathcal{G}_{t+1}^u = \{v_{t+P-L+2}, v_{t+P-L+3}, \dots, v_{t+P+1}\}$  is user  $u$ 's global interaction sequence at  $t+1$ .  $\mathcal{J} = \{j_1, j_2, \dots, j_L\}$  is the negative item sequence containing  $L$  negative items that user has never interact with.  $\mathcal{G}_{t,l}^u = \{v_m, v_{m+1}, \dots, v_{m+l-1}\}$  is a sub-sequence of  $\mathcal{G}_t^u$  and  $m = t + P - L + 1$ .

We share the item embedding and positional embedding between past information encoder and future information encoder, which can bring us two benefits: (1) ensure the consistency of item features between encoders, making the model training stable; (2) reduce the number of trainable parameters for easier model training.

## 2.5 Oracle-Guiding Module

Oracle-Guiding Module is designed to use the future information, as a type of posterior knowledge, to guide the past information to capture the dynamics of user preferences in the future, thereby correcting the deviations in user preference modeling caused by merely using past information. Given two user interaction sequences  $\mathcal{H}_t^u$  and  $\mathcal{G}_t^u$ , we can obtain two item embeddings  $\mathbf{Q}$  and  $\mathbf{R}$  from past and future information encoder respectively. Our goal is to minimize the discrepancy between the predicted embedding of target interaction  $v_t$  from past information encoder, i.e.,  $\mathbf{Q}_L$  (past information) and the predicted embeddings of the item  $v_t$  and the following  $P+1$  items from future information encoder, i.e.,  $\mathbf{R}_{L-P-2+i}$  ( $i = 1, 2, \dots, P+2$ ) (future information), making the former embedding can capture how user preferences change in the future through the latter embeddings. Thus, the loss function of oracle-guiding module is:

$$\mathcal{L}_g = \sum_{i=1}^{P+2} \alpha_i \cdot f(\mathbf{Q}_L, \mathbf{R}_{L-P-2+i}), \quad (13)$$

where  $\alpha_i \in (0, 1]$  is a weight used to distinguish the importance of the  $P+2$  discrepancies. Generally, the guiding should pay more attention to the future information near the target interaction  $v_t$  and pay less attention to the future information far away from  $v_t$ , therefore we use exponential attenuation to model  $\alpha_i$ :

$$\alpha_i = e^{-\gamma(i-1)}, \quad (14)$$

where  $\gamma$  is the attenuation coefficient. Since  $\mathbf{Q}_L$  and  $\mathbf{R}_{L-P-1}$  corresponds to the same item  $v_t$ , the discrepancy  $f(\mathbf{Q}_L, \mathbf{R}_{L-P-1})$  has the highest importance, thus the weight  $\alpha_1$  is 1.  $f(\cdot)$  is a discrepancy measurement function. There are many choices to model  $f(\cdot)$ , we introduce four commonly used choices, including *Kullback-Leibler divergence* (KL divergence), *Jensen-Shannon Divergence* (JS Divergence), *Euclidean distance* and *Cosine distance*. Note that for KL divergence and JS Divergence, we need to use the Softmax function to transform  $\mathbf{Q}$  and  $\mathbf{R}$  to the probabilistic distributions first. We choose KL divergence in Oracle4Rec since it achieves the best performance in empirical studies.

## 2.6 2PTraining: A Tailored Training Strategy

Traditional training, which jointly trains past and future information encoder, is unsuitable for Oracle4Rec since joint training will make the past and future information be extracted simultaneously, weakening the guiding effect of future information. Thus, we tailor an effective two-phase model training strategy 2PTraining, which first trains the future information encoder with the loss  $\mathcal{L}_f$  to obtain future information, then jointly trains past information encoder and oracle-guiding module with the loss  $\mathcal{L}_p + \beta\mathcal{L}_g$  to obtain past information and achieve guiding from future information to past information, to facilitate learning of forward-looking models. We leave the algorithm of 2PTraining in the Appendix due to the space limitation. The experimental results demonstrate the superiority of 2PTraining compared with traditional training strategy.

## 2.7 Inference

In the inference phase, the future information is not available, thus we only use past information encoder to predict the items that users will interact with. Given user  $u$ 's historical interaction sequence  $\mathcal{H}_t^u$ , we first use past information encoder to obtain the predicted item embedding  $\mathbf{Q} \in \mathbb{R}^{L \times d}$ . Since we want to predict which item  $v \in \mathcal{V}$  user  $u$  will interact with after interacting the item  $v_{t-1}$ , we can first calculate the interaction score  $s_{t,v}^u$ , and then the item with the highest score can be seen as the predicted item  $v_t$  that user will be interested to interact with:

$$s_{t,v}^u = \mathbf{Q}_L^T \mathbf{T}_v, \quad v_t = \arg \max_{v \in \mathcal{V}} s_{t,v}^u. \quad (15)$$

## 2.8 Discussions

**2.8.1 Rationality of Adopting Future Information.** Intuitively, the current prediction can be more accurate when taking future uncertainty into account, since future information, as a kind of hindsight observations, can serve as a regularizer to narrow down the representation learning space so as to enhance the accuracy of representation learning and correct the impact of some noisy interactions on user/item representation learning. In deep reinforcement learning, researchers adopt hindsight observation to improve the quality of representation learning [9, 10]. For instance, OPD [9] trained a teacher network with hindsight information, and employed network distillation to facilitate the learning of student network for stock trading. Inspired by this, we design the oracle-guiding module with the hope that future observations can guide and facilitate the current user preference modeling, thereby more accurately predicting user future interactions.

**2.8.2 Differences Between Oracle4Rec and Existing Methods.** Though Oracle4Rec, DualRec and GRec adopt past and future information to model dynamic user preferences, there are five major differences between them, and we leave the detailed comparison in the Appendix due to the space limitation: (1) **The future information extraction manner is different.** DualRec and GRec adopts a right-to-left manner while Oracle4Rec adopts a left-to-right manner. (2) **The future information utilization process is different.** DualRec and GRec learns and utilizes future information simultaneously, whereas Oracle4Rec first learns future information and then utilizes it to guide model training. (3) **The solution to noisy interactions is different.** DualRec and GRec do not consider the impact of

noisy interactions on modeling user preference, while Oracle4Rec proposes a lightweight noise filtering module to deal with noisy interactions. (4) **The training target for future information encoder is different.** DualRec uses the same loss function for both past and future information encoders, while Oracle4Rec redesigns the future information encoder's loss function (Eq.(12)) by considering the guiding role of future information. (5) **The treatment of future information across periods is different.** DualRec treats future information across periods equally while Oracle4Rec treats them selectively by assigning them respective weights when narrowing the gap between future and past information.

**2.8.3 Comparison between Embedding Alignment and Mask & Prediction.** The Mask & Prediction technique is commonly used in the NLP domain [7] to infer the embedding of a masked word based on its contextual information. BERT4Rec [37] inherits this training paradigm and applies it to recommendation algorithms, aiming to infer masked interactions from user interaction sequences. However, we anticipate that current interaction may not always be accurately inferred from future interactions. Therefore, we propose to use the embedding alignment technique, which employs two independent encoders to encode past and future information separately. We use future information to guide model learning on past information. In this training paradigm, our goal is not to infer the current interaction from its context but to regularize the learning of past information using future information, thereby enhancing the consistency of the two types of information. This helps prevent learning spurious correlations, ultimately improving the coherence and accuracy of user preferences modeling.

Moreover, when user interaction sequence exhibits a degree of randomness and noise, it is more challenging for Mask & Prediction technique to infer current interaction from future ones. However, encoding user features to capture future changes and guiding model learning on past features are relatively straightforward.

**Table 1: The statistics of the six real-world datasets.**

	ML100K	ML1M	Beauty	Sports	Toys	Yelp
# Users	943	6,040	22,363	25,598	19,412	30,431
# Items	1,349	3,416	12,101	18,357	11,924	20,033
# Interactions	99,287	999,611	198,052	296,337	167,597	316,354
Density	7.805%	4.485%	0.073%	0.063%	0.072%	0.052%

## 3 EXPERIMENTS

### 3.1 Experimental Setup

We adopt six widely used datasets to comprehensively evaluate the performance of Oracle4Rec: (1) **ML100K and ML1M** (two movie datasets) [13]. (2) **Beauty, Sports and Toys** (three product datasets) [16, 26]. (3) **Yelp** (a business dataset) [1]. To make a fair comparison, for all datasets, we group the interaction records by users, and sort them ascendingly in chronological order. Besides, we follow FMLP-Rec to filter unpopular items and inactive users with fewer than five interactions. We also adopt the same dataset split, which is commonly used in sequential recommendation, as FMLP-Rec, i.e., the last item of each user interaction sequence for

**Table 2: Performance comparison on six datasets. The best result is denoted in bold, the second best result is denoted with an underline. The “★” denotes the statistical significance ( $p < 0.05$ ) of the results of Oracle4Rec compared to the strongest baseline.**

Datasets	Metrics	PopRec	GRU4Rec	Caser	RepeatNet	HGN	CLEA	SASRec	BERT4Rec	GRec	SRGNN	GCSAN	FMLP-Rec	DualRec	Oracle4Rec
ML100K	HR@1	0.0859	0.1737	0.1680	0.1773	0.1513	0.1741	0.1926	0.1786	0.1669	0.1834	0.2040	0.2085	0.2097	<b>0.2257*</b>
	HR@5	0.2397	0.4795	0.4676	0.4732	0.4427	0.4783	0.5005	0.4829	0.4468	0.4806	0.4993	<u>0.5224</u>	0.4993	<b>0.5330*</b>
	NDCG@5	0.1619	0.3300	0.3220	0.3296	0.2971	0.3294	0.3499	0.3349	0.3117	0.3370	0.3581	<u>0.3705</u>	0.3579	<b>0.3843*</b>
	HR@10	0.3669	0.6566	0.6473	0.6431	0.6036	0.6615	0.6653	0.6361	0.6348	0.6585	0.6664	<u>0.6783</u>	0.6670	<b>0.6908*</b>
	NDCG@10	0.2024	0.3872	0.3801	0.3846	0.3514	0.3897	0.4034	0.3846	0.3729	0.3946	0.4121	<u>0.4208</u>	0.4108	<b>0.4353*</b>
	MRR	0.1771	0.3205	0.3144	0.3211	0.2891	0.3211	0.3381	0.3243	0.3098	0.3291	0.3489	<u>0.3562</u>	0.3481	<b>0.3706*</b>
ML1M	HR@1	0.0904	0.3132	0.3119	0.3530	0.2404	0.3203	0.3478	0.3375	0.3448	0.3404	<u>0.3648</u>	0.3541	0.3482	<b>0.3709*</b>
	HR@5	0.3002	0.6576	0.6375	0.6547	0.5644	0.6557	0.6803	0.6632	0.6327	0.6474	0.6694	<u>0.6830</u>	0.6557	<b>0.7106*</b>
	NDCG@5	0.1964	0.4962	0.4848	0.5138	0.4104	0.5128	0.5255	0.5114	0.4978	0.5034	0.5272	<u>0.5294</u>	0.5116	<b>0.5531*</b>
	HR@10	0.4416	0.7754	0.7585	0.7651	0.7087	0.7658	0.7889	0.7716	0.7460	0.7567	0.7773	<u>0.7916</u>	0.7727	<b>0.8128*</b>
	NDCG@10	0.2420	0.5345	0.5236	0.5498	0.4572	0.5486	0.5608	0.5467	0.5337	0.5389	0.5623	<u>0.5648</u>	0.5496	<b>0.5863*</b>
	MRR	0.2038	0.4689	0.4607	0.4923	0.3922	0.4911	0.4982	0.4857	0.4784	0.4811	<u>0.5042</u>	0.5024	0.4901	<b>0.5228*</b>
Beauty	HR@1	0.0678	0.1599	0.1304	0.1613	0.1609	0.1327	0.1794	0.1684	0.1317	0.1777	0.1985	0.2038	<u>0.2182</u>	<b>0.2199*</b>
	HR@5	0.2105	0.3565	0.3054	0.3287	0.3531	0.3310	0.3880	0.3586	0.2868	0.3650	0.3792	0.4089	<u>0.4103</u>	<b>0.4317*</b>
	NDCG@5	0.1391	0.2627	0.2211	0.2482	0.2608	0.2344	0.2886	0.2674	0.2127	0.2753	0.2932	0.3117	<u>0.3192</u>	<b>0.3316*</b>
	HR@10	0.3386	0.4596	0.4072	0.4239	0.4544	0.4425	0.4885	0.4598	0.3826	0.4662	0.4707	<u>0.5079</u>	0.5029	<b>0.5295*</b>
	NDCG@10	0.1803	0.2960	0.2539	0.2789	0.2934	0.2704	0.3211	0.3000	0.2435	0.3078	0.3227	0.3436	<u>0.3490</u>	<b>0.3632*</b>
	MRR	0.1558	0.2633	0.2264	0.2524	0.2616	0.2378	0.2865	0.2689	0.2209	0.2768	0.2940	0.3092	<u>0.3178</u>	<b>0.3278*</b>
Sports	HR@1	0.0763	0.1291	0.1024	0.1325	0.1444	0.1254	0.1503	0.1405	0.1055	0.1473	0.1674	0.1699	<u>0.1831</u>	<b>0.1861*</b>
	HR@5	0.2293	0.3397	0.2834	0.3208	0.3484	0.3397	0.3672	0.3466	0.2786	0.3498	0.3702	0.3915	<u>0.4031</u>	<b>0.4191*</b>
	NDCG@5	0.1538	0.2372	0.1947	0.2290	0.2492	0.2327	0.2620	0.2463	0.1937	0.2517	0.2721	0.2846	<u>0.2981</u>	<b>0.3070*</b>
	HR@10	0.3423	0.4635	0.4061	0.4381	0.4698	0.4623	0.4964	0.4730	0.3948	0.4738	0.4909	0.5137	<u>0.5267</u>	<b>0.5431*</b>
	NDCG@10	0.1902	0.2772	0.2342	0.2668	0.2883	0.2750	0.3037	0.2870	0.2315	0.2916	0.3111	0.3240	<u>0.3380</u>	<b>0.3471*</b>
	MRR	0.1660	0.2391	0.2028	0.2342	0.2522	0.2504	0.2632	0.2496	0.2027	0.2552	0.2743	0.2833	<u>0.2980</u>	<b>0.3037*</b>
Toys	HR@1	0.0585	0.1481	0.1114	0.1370	0.1541	0.1245	0.1799	0.1504	0.1069	0.1682	0.1995	0.1926	<b>0.2170</b>	<u>0.2152</u>
	HR@5	0.1977	0.3456	0.2944	0.3007	0.3430	0.3278	0.3652	0.3446	0.2687	0.3572	0.3787	0.4007	<u>0.4043</u>	<b>0.4233*</b>
	NDCG@5	0.1286	0.2505	0.2054	0.2213	0.2516	0.2286	0.2766	0.2504	0.1891	0.2661	0.2926	0.3015	<u>0.3149</u>	<b>0.3246*</b>
	HR@10	0.3008	0.4532	0.4052	0.4007	0.4485	0.4421	0.4609	0.4567	0.3694	0.4605	0.4740	<u>0.5032</u>	0.4999	<b>0.5222*</b>
	NDCG@10	0.1618	0.2852	0.2411	0.2535	0.2857	0.2643	0.3075	0.2866	0.2219	0.2995	0.3234	0.3346	<u>0.3458</u>	<b>0.3565*</b>
	MRR	0.1430	0.2517	0.2107	0.2279	0.2542	0.2313	0.2778	0.2531	0.1983	0.2680	0.2942	0.2994	<u>0.3153</u>	<b>0.3215*</b>
Yelp	HR@1	0.0801	0.2109	0.1858	0.2344	0.2337	0.2085	0.2276	0.2446	0.1711	0.2400	0.2561	0.2704	<u>0.2817</u>	<b>0.2970*</b>
	HR@5	0.2415	0.5772	0.5170	0.5360	0.5663	0.5678	0.5856	0.5848	0.4743	0.5772	0.5932	0.6231	<u>0.6272</u>	<b>0.6499*</b>
	NDCG@5	0.1622	0.3998	0.3560	0.3897	0.4057	0.3842	0.4117	0.4203	0.3251	0.4147	0.4301	0.4538	<u>0.4610</u>	<b>0.4812*</b>
	HR@10	0.3609	0.7474	0.6866	0.6950	0.7306	0.7488	0.7651	0.7514	0.6567	0.7398	0.7557	0.7777	<u>0.7839</u>	<b>0.8000*</b>
	NDCG@10	0.2007	0.4552	0.4110	0.4411	0.4589	0.4465	0.4699	0.4744	0.3837	0.4675	0.4829	0.5040	<u>0.5119</u>	<b>0.5300*</b>
	MRR	0.1740	0.3770	0.3409	0.3778	0.3885	0.3669	0.3915	0.4017	0.3176	0.3966	0.4113	0.4299	<u>0.4393</u>	<b>0.4562*</b>

test, the penultimate item for validation, and all remaining items for training. Table 1 describes the statistics of the six datasets.

We compare Oracle4Rec with 13 sequential methods. (1) **PopRec**; (2) **GRU4Rec** [17]; (3) **Caser** [39]; (4) **HGN** [25]; (5) **RepeatNet** [31]; (6) **CLEA** [29]; (7) **SASRec** [21]; (8) **BERT4Rec** [37]; (9) **GRec** [53]; (10) **SRGNN** [45]; (11) **GCSAN** [50]; (12) **FMLP-Rec** [55]; and (13) **DualRec** [54]. For all methods, BERT4Rec, GRec and DualRec can access the future information, while other methods can merely use past information for user preference modeling.

We compare Oracle4Rec with other methods in top-K recommendation task with three popular ranking metrics: (1) **Hit Ratio (HR)**; (2) **Normalized Discounted Cumulative Gain (NDCG)**; and (3) **Mean Reciprocal Rank (MRR)**. For the former two metrics, we report their results when  $K=1, 5, 10$ , i.e.,  $HR@\{1, 5, 10\}$  and  $NDCG@\{1, 5, 10\}$ , and we omit  $NDCG@1$  since it is equal to  $HR@1$ . For a fair comparison, we follow the same negative sampling strategy as FMLP-Rec to pair the ground-truth item with 99 randomly sampled negative items that the user has not interacted with.

We adopt Adam optimizer [22] to optimize the model with learning rates  $\eta_1 = \eta_2 = 0.001$ , we set the maximum sequence length  $L$  to 50 and the number of future interactions  $P$  in user global interaction sequence to 10. For all datasets, we tune hidden size  $d$

from [32, 64, 128, 256], attenuation coefficient  $\gamma$  from 0.005 to 1.0, frequency quantile  $q$  from 0.3 to 1.0 and regularization coefficient  $\beta$  from [0.005, 0.01, 0.05]. Dropout probabilities are fixed to 0.5. We also tune the layer number of noise filtering module from 1 to 3 and the layer number of causal self-attention module from 1 to 5. The hyper-parameters of all baselines are carefully tuned according to their papers. Due to the space limitation, we leave the detailed setup in the Appendix.

### 3.2 Performance Comparison

Table 2 shows performance comparison of all methods. Following FMLP-Rec, we also report the full-ranking results, i.e., ranking the ground-truth item with all candidate items, in the Appendix. From the results, we have the following observations:

1. *Transformer-based methods (i.e., SASRec and BERT4Rec) achieve better accuracy than RNN-based methods (i.e., GRU4Rec and HGN) and CNN-based methods (i.e., Caser).* The main reason is that the self-attention mechanism in Transformer-based methods has the largest receptive field compared with RNNs and CNNs, so it can capture more information from user interaction sequence, which can make user preference modeling more precise.

2. *GNN-based methods (i.e., SRGNN and GCSAN) have comparable or better performance than Transformer-based methods.* This is because the GNN has a powerful structure feature extraction ability [51], and can capture the transition relationship between items, so as to achieve more accurate recommendation results.

3. *FMLP-Rec and DualRec achieve better accuracy than all other baselines.* This is because the former has strong noise reduction capability to filter the noise in the user interactions, and the latter can utilize both past information and future information from user interactions, making them model user preference more accurate.

4. *Oracle4Rec consistently outperforms all compared methods on all datasets, demonstrating its superiority.* Compared with methods only using past information, Oracle4Rec effectively leverages past and future information in training, making the dynamic user preference modeling more accurate. Compared with BERT4Rec, GRec and DualRec, Oracle4Rec adopts an oracle-guiding module and a 2PTraining to leverage future information in a more effective way, thus achieving better performance. Generally, using future information can reduce the error of predicting user’s interested items during training, and thus will make more accurate recommendation. This can be drawn by comparing the losses, which is placed in the Appendix due to the space limitation, when Oracle4Rec uses past information only or past and future information together.

### 3.3 Ablation Study

We conduct the extensive ablation study to comprehensively analyze the impact of each component on the performance of Oracle4Rec. Table 3 shows the results. Due to the space limitation, the results of other metrics are placed in the Appendix.

1. Comparing setting (1) and (11), we can find that without the noise filtering module, the performance of Oracle4Rec significantly decreases on all cases. This is because user interactions are inevitably noisy. By reducing the noise in the interaction sequence using low-pass filter, Oracle4Rec can more accurately model user preferences, thus making more accurate recommendations. Comparing setting (2) and (11), the results show that low-pass filter (setting (11)) achieves better results than learnable filter, which we attribute to the additional parameters of learnable filter increase the difficulty of model training, thus reducing the performance.

2. Comparing setting (3) and (11), we can find that Oracle4Rec obtains better accuracy than Oracle4Rec w/o Future Information Encoder, which shows the importance of oracle information guiding in model training, i.e., punish deviations from user future preferences. Moreover, merely using future information encoder can also lead to sub-optimal performance by comparing setting (9) and (11). Therefore, we can conclude that guiding from future information to past information can make user preference modeling more reliable.

3. The attenuation in the oracle-guiding module is helpful to improve the model performance by comparing setting (4) and (11). This is because with attenuation, Oracle4Rec can selectively minimize the discrepancies between past and future information from different time, thus capturing the evolution of user preferences and achieving better performance.

4. We can find that 2PTraining is beneficial to improve the performance of Oracle4Rec by comparing setting (5) and (11). Traditional training strategy jointly trains past and future information encoders,

which is unable to sufficiently leverage past information and future information. However, 2PTraining first optimizes future information encoder then optimizes the past information encoder, making the model able to fully leverage these two information.

5. Comparing setting (6)–(8) and (11), we can find that the performance of Oracle4Rec varies with different discrepancy measurement methods. When equipped with KL divergence, Oracle4Rec achieves the best results. Thus, we adopt KL divergence as the discrepancy measurement method in Oracle4Rec.

6. The L2R-style (left-to-right-style) information extraction manner is superior than the R2L-style (right-to-left-style) information extraction manner by comparing setting (10) and (11), which demonstrates the rationality of our left-to-right-style. Generally, the left-to-right-style is more intuitive and aligns with the chronological order of interactions, making the extracted information more sufficient and the modeled user preference more accurate.

### 3.4 Future Information Encoder Analysis

To explore how future information encoder can facilitate the learning process of Oracle4Rec, or in other words, what Oracle4Rec will learn when equipped with future information encoder, we analyze the consistency of user real preference and user predicted preference on ML100K and ML1M. The former preference is calculated from the categories of real items that user interacts with in the inference phase, which is denoted as  $p$ , and the latter preference is calculated from the categories of top-10 items predicted by Oracle4Rec w/o future information encoder and Oracle4Rec w/ future information encoder, denoted as  $q^{(1)}$  and  $q^{(2)}$ . By comparing the KL divergence  $\text{KL}(p, q^{(1)})$  and  $\text{KL}(p, q^{(2)})$ , we can determine whether future information encoder benefits the learning process of Oracle4Rec. We leave the details about the calculation of preference and KL divergence in the Appendix due to the space limitation.

Table 4 shows the experimental results. We can observe that the KL divergence is reduced when Oracle4Rec is equipped with future information encoder, which demonstrates that future information encoder can make predicted items more similar to items that user will interact with in terms of item categories and is beneficial to the leaning process of Oracle4Rec. Generally, future information encoder can help infer “oracle” user preference in the future, and with the help of oracle-guiding module, Oracle4Rec can recognize how user preference will change from past to future, which cannot be achieved by merely using past information encoder, thus making user dynamic preference modeling more accurate.

### 3.5 Generality Analysis

We apply the oracle-guiding module and 2PTraining of Oracle4Rec into four sequential methods GRU4Rec, SRGNN, GCSAN and FMLP-Rec and evaluate their performance afterwards to further analyze the generality of Oracle4Rec. Table 5 shows the results, the results of other metrics are placed in the Appendix. We can find that the performance of all methods has a significant improvement, which shows Oracle4Rec has high generality. The results also show that leveraging future information can improve the performance of sequential methods, since future information as a type of posterior knowledge contains the evolution of user preferences in the future,



**Table 3: Ablation study of Oracle4Rec on ML100K, ML1M and Beauty datasets. Bold face indicates the highest performance.**

	ML100K			ML1M			Beauty		
	HR@1	NDCG@5	MRR	HR@1	NDCG@5	MRR	HR@1	NDCG@5	MRR
(1) Oracle4Rec w/o Noise Filtering Module	0.1909	0.3554	0.3443	0.3658	0.5495	0.5189	0.2062	0.3104	0.3090
(2) Oracle4Rec w/ Learnable Filter	0.2042	0.3637	0.3513	0.3523	0.5262	0.4995	0.2087	0.3121	0.3110
(3) Oracle4Rec w/o Future Information Encoder	0.1782	0.3391	0.3277	0.3477	0.5275	0.4993	0.1878	0.3010	0.2975
(4) Oracle4Rec w/o Attenuation of Discrepancies	0.1983	0.3604	0.3491	0.3347	0.5225	0.4924	0.2123	0.3255	0.3217
(5) Oracle4Rec w/ Traditional Training Strategy	0.2087	0.3670	0.3566	0.3497	0.5302	0.5016	0.2041	0.3164	0.3129
(6) Oracle4Rec w/ JS Diverge	0.1968	0.3583	0.3456	0.3555	0.5399	0.5095	0.2160	0.3257	0.3222
(7) Oracle4Rec w/ Euclidean Distance	0.1987	0.3622	0.3491	0.3630	0.5445	0.5149	0.2071	0.3194	0.3158
(8) Oracle4Rec w/ Cosine Distance	0.2070	0.3712	0.3572	0.3608	0.5439	0.5137	0.2078	0.3211	0.3172
(9) Oracle4Rec Future Information Encoder	0.1911	0.3407	0.3328	0.3283	0.5063	0.4798	0.1986	0.3017	0.3004
(10) Oracle4Rec w/ R2L-style	0.2078	0.3659	0.3549	0.3389	0.5157	0.4892	0.2071	0.3174	0.3148
(11) Oracle4Rec (w/ KL Divergence & L2R-style)	<b>0.2257</b>	<b>0.3843</b>	<b>0.3706</b>	<b>0.3709</b>	<b>0.5531</b>	<b>0.5228</b>	<b>0.2199</b>	<b>0.3316</b>	<b>0.3278</b>

**Table 4: The KL divergence between user real and predicted preference under different settings on ML100K and ML1M.**

	ML100K	ML1M
Oracle4Rec w/o Future Information Encoder	0.0085	0.0069
Oracle4Rec w/ Future Information Encoder	0.0070 (+17.6%)	0.0062 (+10.1%)

**Table 5: Generality analysis of four sequential methods. “XYZ+” means applying Oracle4Rec into the corresponding method XYZ. Bold face indicates better performance in that group. The “★” denotes the statistical significance ( $p < 0.05$ ) of the results of XYZ+ compared to XYZ.**

	ML100K			ML1M		
	HR@1	NDCG@5	MRR	HR@1	NDCG@5	MRR
GRU4Rec	0.1737	0.3300	0.3205	0.3132	0.4962	0.4689
GRU4Rec+	<b>0.2059★</b>	<b>0.3642★</b>	<b>0.3532★</b>	<b>0.3446★</b>	<b>0.5244★</b>	<b>0.4964★</b>
SRGNN	0.1834	0.3370	0.3291	0.3404	0.5034	0.4811
SRGNN+	<b>0.1985★</b>	<b>0.3574★</b>	<b>0.3466★</b>	<b>0.3578★</b>	<b>0.5244★</b>	<b>0.5006★</b>
GCSAN	0.2040	0.3581	0.3489	0.3648	0.5272	0.5042
GCSAN+	<b>0.2159★</b>	<b>0.3741★</b>	<b>0.3613★</b>	<b>0.3846★</b>	<b>0.5462★</b>	<b>0.5231★</b>
FMLP-Rec	0.2085	0.3705	0.3562	0.3541	0.5294	0.5024
FMLP-Rec+	<b>0.2178★</b>	<b>0.3788★</b>	<b>0.3658★</b>	<b>0.3657★</b>	<b>0.5396★</b>	<b>0.5125★</b>

which can address the deviations caused by only using past information to model dynamic user preferences as GRU4Rec, SRGNN, GCSAN and FMLP-Rec do, thus achieving superior performance.

## 4 RELATED WORK

Sequential recommendation methods model dynamic user preferences according to his/her historical interactions [25, 29, 31]. Several techniques have been applied in this line of work, from the earliest Markov Chain [15, 32], to the later Recurrent Neural Networks [17, 44], Convolutional Neural Networks [39, 41], Self-Attention Mechanism [21, 37] and Graph Neural Networks [45, 50]. FPMC [32] is a Markov Chain-based sequential method that combines Markov chain and matrix factorization to model user preferences. Fossil [15] fuses the similarity-based method with Markov

Chain to characterize users in terms of both preferences and the strength of sequential behavior, addressing the sparsity issues and the long-tailed distribution of datasets. Due to the vigorous development of deep learning, the performance of sequential recommendation methods has been significantly improved. GRU4Rec [17] adopts GRU [6] to model the dynamic user preferences. 3D-CNN [41] uses a convolutional neural network to capture sequential patterns and model user preference. SASRec [21] uses a self-attention mechanism [42] to capture user long-term preferences. SRGNN [45] and GCSAN [50] are GNN-based sequential methods that use graph neural networks [23, 33] and self-attention mechanism to model user preferences. FMLP-Rec [55] adopts an all-MLP architecture to model dynamic user preferences. Note that contrastive learning based sequential methods are another line of works, they are different from ours, thus we exclude them.

However, the above methods only use the past information to make recommendation, making the model performance sub-optimal. Future information is also available during training and is beneficial to model user preferences. BERT4Rec [37], GRec [53] and DualRec [54] are three sequential methods that integrates past and future information to model dynamic user preferences. Though Oracle4Rec also leverages future information to model dynamic user preferences, there are several major differences between existing methods and Oracle4Rec, as discussed in Section 2.8.2.

## 5 CONCLUSION

We propose an oracle-guided dynamic user preference modeling method for sequential recommendation that uses future information to guide model training on past information, so as to better model dynamic user preferences. Oracle4Rec first extracts past and future information through two separate encoders, then learns forward-looking models through an oracle-guiding module and a tailored model training strategy. Extensive experimental results demonstrate the superiority and generality of Oracle4Rec.

## ACKNOWLEDGMENTS

This research was supported by National Natural Science Foundation of China (NSFC) under the Grant No. 62172106, 61932007, and 62372113.



## ETHICAL CONSIDERATIONS

To the best of our knowledge, this work does not bring new negative societal impacts, including fairness, privacy, security, safety, misuse of the technology by malicious actors, as well as possible harms that could arise even when the technology is being used as intended and functioning correctly.

## REFERENCES

- [1] Nabihha Asghar. 2016. Yelp dataset challenge: Review rating prediction. *arXiv preprint arXiv:1605.05362* (2016).
- [2] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E Hinton. 2016. Layer normalization. *arXiv preprint arXiv:1607.06450* (2016).
- [3] Soma Bandyopadhyay and SS Thakur. 2020. Product prediction and recommendation in e-commerce using collaborative filtering and artificial neural networks: A hybrid approach. In *Intelligent Computing Paradigm: Recent Trends*. Springer, 59–67.
- [4] Rianne van den Berg, Thomas N Kipf, and Max Welling. 2017. Graph convolutional matrix completion. *arXiv preprint arXiv:1706.02263* (2017).
- [5] Heng-Tze Cheng, Levent Koc, Jeremiah Harmsen, Tal Shaked, Tushar Chandra, Hrishikesh Aradhye, Glen Anderson, Greg Corrado, Wei Chai, Mustafa Ipsir, et al. 2016. Wide & deep learning for recommender systems. In *Proceedings of the 1st workshop on deep learning for recommender systems*. 7–10.
- [6] Kyunghyun Cho, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. 2014. Learning phrase representations using RNN encoder-decoder for statistical machine translation. *arXiv preprint arXiv:1406.1078* (2014).
- [7] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805* (2018).
- [8] Pierre Duhamel and Martin Vetterli. 1990. Fast Fourier transforms: a tutorial review and a state of the art. *Signal processing* 19, 4 (1990), 259–299.
- [9] Yuchen Fang, Kan Ren, Weiqing Liu, Dong Zhou, Weinan Zhang, Jiang Bian, Yong Yu, and Tie-Yan Liu. 2021. Universal trading for order execution with oracle policy distillation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 35. 107–115.
- [10] Arthur Guez, Fabio Viola, Théophane Weber, Lars Buesing, Steven Kapturowski, Doina Precup, David Silver, and Nicolas Heess. 2020. Value-driven hindsight modelling. *Advances in Neural Information Processing Systems* 33 (2020), 12499–12509.
- [11] Xudong Guo, Xun Guo, and Yan Lu. 2021. Ssan: Separable self-attention network for video representation learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 12618–12627.
- [12] Meenu Gupta, Aditya Thakkar, Vishal Gupta, Dhruv Pratap Singh Rathore, et al. 2020. Movie recommender system using collaborative filtering. In *2020 International Conference on Electronics and Sustainable Communication Systems (ICESC)*. IEEE, 415–420.
- [13] F Maxwell Harper and Joseph A Konstan. 2015. The movielens datasets: History and context. *Acm transactions on interactive intelligent systems (tiis)* 5, 4 (2015), 1–19.
- [14] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016. Deep residual learning for image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*. 770–778.
- [15] Ruining He and Julian McAuley. 2016. Fusing similarity models with markov chains for sparse sequential recommendation. In *2016 IEEE 16th international conference on data mining (ICDM)*. IEEE, 191–200.
- [16] Ruining He and Julian McAuley. 2016. Ups and downs: Modeling the visual evolution of fashion trends with one-class collaborative filtering. In *proceedings of the 25th international conference on world wide web*. 507–517.
- [17] Balázs Hidasi, Alexandros Karatzoglou, Linas Baltrunas, and Domonkos Tikk. 2015. Session-based recommendations with recurrent neural networks. *arXiv preprint arXiv:1511.06939* (2015).
- [18] Sepp Hochreiter and Jürgen Schmidhuber. 1997. Long short-term memory. *Neural computation* 9, 8 (1997), 1735–1780.
- [19] Katja Hofmann, Fritz Behr, and Filip Radlinski. 2012. On caption bias in interleaving experiments. In *Proceedings of the 21st ACM international conference on Information and knowledge management*. 115–124.
- [20] Rolf Jagerman, Harrie Oosterhuis, and Maarten de Rijke. 2019. To model or to intervene: A comparison of counterfactual and online learning to rank from user interactions. In *Proceedings of the 42nd international ACM SIGIR conference on research and development in information retrieval*. 15–24.
- [21] Wang-Cheng Kang and Julian McAuley. 2018. Self-attentive sequential recommendation. In *2018 IEEE international conference on data mining (ICDM)*. IEEE, 197–206.
- [22] Diederik P Kingma and Jimmy Ba. 2014. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980* (2014).
- [23] Thomas N Kipf and Max Welling. 2016. Semi-supervised classification with graph convolutional networks. *arXiv preprint arXiv:1609.02907* (2016).
- [24] Dongsheng Li, Jianxun Lian, Le Zhang, Kan Ren, Tun Lu, Tao Wu, and Xing Xie. 2024. *Recommender Systems: Frontiers and Practices*. Springer Nature.
- [25] Chen Ma, Peng Kang, and Xue Liu. 2019. Hierarchical gating networks for sequential recommendation. In *Proceedings of the 25th ACM SIGKDD international conference on knowledge discovery & data mining*. 825–833.
- [26] Julian McAuley, Christopher Targett, Qinfeng Shi, and Anton Van Den Hengel. 2015. Image-based recommendations on styles and substitutes. In *Proceedings of the 38th international ACM SIGIR conference on research and development in information retrieval*. 43–52.
- [27] Andriy Mnih and Russ R Salakhutdinov. 2007. Probabilistic matrix factorization. *Advances in neural information processing systems* 20 (2007).
- [28] Witriana Endah Pangesti, Rachmat Suryadithia, Muhammad Faisal, Bilal Abdul Wahid, Arman Syah Putra, et al. 2021. Collaborative Filtering Based Recommender Systems For Marketplace Applications. *International Journal of Educational Research and Social Sciences (IJERSC)* 2, 5 (2021), 1201–1209.
- [29] Yuqi Qin, Pengfei Wang, and Chenliang Li. 2021. The world is binary: Contrastive learning for denoising next basket recommendation. In *Proceedings of the 44th International ACM SIGIR Conference on Research and Development in Information Retrieval*. 859–868.
- [30] K RaviKanth, K ChandraShekar, K Sreekanth, and P Santhosh Kumar. 2019. Recommendation system for e-commerce by memory based and model based collaborative filtering. In *International Conference on Soft Computing and Pattern Recognition*. Springer, 123–129.
- [31] Pengjie Ren, Zhumin Chen, Jing Li, Zhaochun Ren, Jun Ma, and Maarten De Rijke. 2019. Repeatnet: A repeat aware neural recommendation machine for session-based recommendation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 33. 4806–4813.
- [32] Steffen Rendle, Christoph Freudenthaler, and Lars Schmidt-Thieme. 2010. Factorizing personalized markov chains for next-basket recommendation. In *Proceedings of the 19th international conference on World wide web*. 811–820.
- [33] Franco Scarselli, Marco Gori, Ah Chung Tsoi, Markus Hagenbuchner, and Gabriele Monfardini. 2008. The graph neural network model. *IEEE transactions on neural networks* 20, 1 (2008), 61–80.
- [34] Suvasish Sedhain, Aditya Krishna Menon, Scott Sanner, and Lexing Xie. 2015. Autorec: Autoencoders meet collaborative filtering. In *Proceedings of the 24th international conference on World Wide Web*. 111–112.
- [35] Mahesh Kumar Singh and Om Prakash Rishi. 2020. Event driven recommendation system for E-commerce using knowledge based collaborative filtering technique. *Scalable Computing: Practice and Experience* 21, 3 (2020), 369–378.
- [36] Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. 2014. Dropout: a simple way to prevent neural networks from overfitting. *The journal of machine learning research* 15, 1 (2014), 1929–1958.
- [37] Fei Sun, Jun Liu, Jian Wu, Changhua Pei, Xiao Lin, Wenwu Ou, and Peng Jiang. 2019. BERT4Rec: Sequential recommendation with bidirectional encoder representations from transformer. In *Proceedings of the 28th ACM international conference on information and knowledge management*. 1441–1450.
- [38] Gongbo Tang, Mathias Müller, Annette Rios, and Rico Sennrich. 2018. Why self-attention? a targeted evaluation of neural machine translation architectures. *arXiv preprint arXiv:1808.08946* (2018).
- [39] Jiayi Tang and Ke Wang. 2018. Personalized top-n sequential recommendation via convolutional sequence embedding. In *Proceedings of the eleventh ACM international conference on web search and data mining*. 565–573.
- [40] Urvish Thakker, Ruhi Patel, and Manan Shah. 2021. A comprehensive analysis on movie recommendation system employing collaborative filtering. *Multimedia Tools and Applications* 80, 19 (2021), 28647–28672.
- [41] Trinh Xuan Tuan and Tu Minh Phuong. 2017. 3D convolutional networks for session-based recommendation with content features. In *Proceedings of the eleventh ACM conference on recommender systems*. 138–146.
- [42] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in neural information processing systems* 30 (2017).
- [43] Wenjie Wang, Fuli Feng, Xiangnan He, Liqiang Nie, and Tat-Seng Chua. 2021. Denoising implicit feedback for recommendation. In *Proceedings of the 14th ACM international conference on web search and data mining*. 373–381.
- [44] Chao-Yuan Wu, Amr Ahmed, Alex Beutel, Alexander J Smola, and How Jing. 2017. Recurrent recommender networks. In *Proceedings of the tenth ACM international conference on web search and data mining*. 495–503.
- [45] Shu Wu, Yuyuan Tang, Yanqiao Zhu, Liang Wang, Xing Xie, and Tieniu Tan. 2019. Session-based recommendation with graph neural networks. In *Proceedings of the AAAI conference on artificial intelligence*, Vol. 33. 346–353.
- [46] Jiafeng Xia, Dongsheng Li, Hansu Gu, Jiahao Liu, Tun Lu, and Ning Gu. 2022. FIRE: Fast incremental recommendation with graph signal processing. In *Proceedings of the ACM Web Conference 2022*. 2360–2369.
- [47] Jiafeng Xia, Dongsheng Li, Hansu Gu, Tun Lu, Peng Zhang, and Ning Gu. 2021. Incremental graph convolutional network for collaborative filtering. In *Proceedings*

- of the 30th ACM International Conference on Information & Knowledge Management. 2170–2179.
- [48] Jiafeng Xia, Dongsheng Li, Hansu Gu, Tun Lu, Peng Zhang, Li Shang, and Ning Gu. 2024. Hierarchical Graph Signal Processing for Collaborative Filtering. In *Proceedings of the ACM on Web Conference 2024*. 3229–3240.
  - [49] Jiafeng Xia, Dongsheng Li, Hansu Gu, Tun Lu, Peng Zhang, Li Shang, and Ning Gu. 2024. Neural Kalman Filtering for Robust Temporal Recommendation. In *Proceedings of the 17th ACM International Conference on Web Search and Data Mining*. 836–845.
  - [50] Chengfeng Xu, Pengpeng Zhao, Yanchi Liu, Victor S Sheng, Jiajie Xu, Fuzhen Zhuang, Junhua Fang, and Xiaofang Zhou. 2019. Graph Contextualized Self-Attention Network for Session-based Recommendation.. In *IJCAI*, Vol. 19. 3940–3946.
  - [51] Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. 2018. How powerful are graph neural networks? *arXiv preprint arXiv:1810.00826* (2018).
  - [52] Wenhui Yu and Zheng Qin. 2020. Graph convolutional network for recommendation with low-pass collaborative filters. In *International Conference on Machine Learning*. PMLR, 10936–10945.
  - [53] Fajie Yuan, Xiangnan He, Haochuan Jiang, Guibing Guo, Jian Xiong, Zhezha Xu, and Yilin Xiong. 2020. Future data helps training: Modeling future contexts for session-based recommendation. In *Proceedings of The Web Conference 2020*. 303–313.
  - [54] Hengyu Zhang, Enming Yuan, Wei Guo, Zhicheng He, Jiarui Qin, Huifeng Guo, Bo Chen, Xiu Li, and Ruiming Tang. 2022. Disentangling Past-Future Modeling in Sequential Recommendation via Dual Networks. In *Proceedings of the 31st ACM International Conference on Information & Knowledge Management*. 2549–2558.
  - [55] Kun Zhou, Hui Yu, Wayne Xin Zhao, and Ji-Rong Wen. 2022. Filter-enhanced MLP is all you need for sequential recommendation. In *Proceedings of the ACM Web Conference 2022*. 2388–2399.